

Web Searching Mining Project2 – Lemur(Indri)

student ID:110753201 Class:資科碩一 Name: 曹昱維

Introduction

- This project is asked to use Lemur or Indri toolkits to implement several different retrieval methods.
- Corpus : WT2g. This corpus contains Web documents with a 2GB corpus.
- Install indri 5.2
 - Install script:

```
echo "deb http://dk.archive.ubuntu.com/ubuntu/ trusty main universe" \
>> /etc/apt/sources.list \
apt update && apt upgrade -y && \
apt install -y build-essential make gcc-4.4 g++-4.4 zlib1g-dev && \
update-alternatives --install /usr/bin/g++ g++ /usr/bin/g++-9 1 && \
update-alternatives --install /usr/bin/gcc gcc /usr/bin/gcc-4.4 2 && \
update-alternatives --install /usr/bin/g++ g++ /usr/bin/g++-4.4 2 && \
update-alternatives --install /usr/bin/gcc gcc /usr/bin/gcc-9 1 && \
tar -xzf /root/indri-5.20.tar.gz -C /root && \
cd /root/indri-5.20 && ./configure && make && make install
```

- The main process consist of three steps:
 1. Buidl Index
 - build two type index of corpus, **index with porter stemming and without stemming.**
 2. Run Query
 - Run query by specified method, for example dirichlet smoothing.
 - We have **four method** to test, Okapi, Language Model with laplace smooth, Language Model with Jelinek-Mercer smoothing, method proposed by student.
 3. evaluation
 - We evaluate 2*4 retrieved result files(two type stemming and 4 methods).

Related work

- The follwing description of How to set up parameters files of IndriBuildIndex , IndriRunQuery refers to
 - <https://sourceforge.net/p/lemur/wiki/Home/> (<https://sourceforge.net/p/lemur/wiki/Home/>)
 - <http://www.cs.cmu.edu/~lemur/1.1/api.html> (<http://www.cs.cmu.edu/~lemur/1.1/api.html>)
 - <https://sourceforge.net/projects/lemur/files/lemur/indri-5.20/indri-5.20-doc.zip/download> (<https://sourceforge.net/projects/lemur/files/lemur/indri-5.20/indri-5.20-doc.zip/download>)
- **IndriBuildIndex** <index_parameter_file>
 - Indri 5.2 provide a useful program called IndriBuildIndex , which build index file of corpus, you can set several options, such as *stem method*, *stop word* etc. in the <index_parameter files>
 - ref: <https://sourceforge.net/p/lemur/wiki/IndriBuildIndex%20Parameters/> (<https://sourceforge.net/p/lemur/wiki/IndriBuildIndex%20Parameters/>)
 - An Example of the IndriBuildIndex **parameter file** used in Project2 of WSM is shown in Figure 1

```
1 <parameters>
2   <index>index/porter_index</index>
3   <corpus>
4     <path>WT2G</path>
5     <class>trecweb</class>
6   </corpus>
7   <memory>1000m</memory>
8   <stemmer>
9     <name>porter</name>
10  </stemmer>
11  <stopper>
12    <word>a</word>
13    <word>about</word>
14    .
15    .
16    .
17    <word>yourselves</word>
18  </stopper>
19 </parameters>
```

Fig.1

- **IndriRunQuery** [query_parameter_file] [query_data] > result file
 - Indri 5.2 provide a useful program called **IndriRunQuery** , which run queries on corpus. The queries should sotred in [query_data] . You can set several options such as *retrieve relvance model*, *smoothing mehtod* etc. in the [query_parameter files] .
 - ref: <https://sourceforge.net/p/lemur/wiki/IndriRunQuery/>
(<https://sourceforge.net/p/lemur/wiki/IndriRunQuery/>)
 - An Example of the **IndriRunQuery** **parameter file** used in Project2 of WSM is shown in Figure 2

```

1 <parameters>
2   <index>index/porter_index</index>
3   <count>1000</count>
4   <baseline>okapi,k1:1.5,b:0.75,k3:1.5</baseline>
5   <trecFormat>true</trecFormat>
6   <queryOffset>401</queryOffset>
7   <runID>Exp</runID>
8 </parameters>

```

Fig.2

- An Example of the **IndriRunQuery** **queries data** used in Project2 of WSM is shown in Figure 3

```

1 <parameters>
2   <query>
3     foreign minorities Germany
4   </query>
5   .
6   .
7   .
8   <query>
9     antibiotics ineffectiveness
10  </query>
11  <query>
12    King Hussein peace
13  </query>
14 </parameters>

```

Fig.3

Run Query

- **Overview of how to run queries**
 - We need to test 4 models: *Okapi TF* , *language model with Laplace smoothing*, *Language modeling with Jelinek-Mercer smoothing*, *improve one of the above three IR models*.
 - We have two index files of coupus: *without stemming*. *with stemming*. In this porject, I choose **Porter setmmmer**.
 - Therefore, with above description, we will get **2 index files** and **8 retrieved results** as shown in Figure 4 and Figure5 respectively.

```

root@de5f2bba8c50 ~/wsm_proj2
# ls index
no_stem_index  porter_index

```

Fig.4

```

ywt01@ywt01-15Z90N-V-AR53C2 ~/Documents/codes/nccu_cs_hw/WebSearchMiningHW/Pro
ject2 (main)
$ ls codes/ret results
ret_no_stem_JM.txt          ret_porter_JM.txt
ret_no_stem_laplace_log.txt ret_porter_laplace_log.txt
ret_no_stem_okapi_k15b075k15.txt ret_porter_okapi_k15b075k15.txt
ret_no_stem_okapi_k2b075k0.txt ret_porter_okapi_k2b075k0.txt

```

Fig.5

- **Preprocess query files**
 - We used title field of original queries file(*topics.401-450.txt*) as input queries to **IndriRunQuery**
 - An example of original queries file(*topics.401-450.txt*) as shown in Figure 6.

```

1 <top>
2
3 <num> Number: 401
4 <title> foreign minorities, Germany
5
6 <desc> Description:
7 What language and cultural differences impede the integration
8 of foreign minorities in Germany?
9
10 <narr> Narrative:
11 A relevant document will focus on the causes of the lack of
12 integration in a significant way; that is, the mere mention of
13 immigration difficulties is not relevant. Documents that discuss
14 immigration problems unrelated to Germany are also not relevant.
15
16 </top>

```

Fig.6

- An example of after-processed query as shown in Figure 7.

```

1 <parameters>
2 <query>
3   foreign minorities Germany
4 </query>

```

Fig.7

- Preprocessing linux command:

```

cat topics.401-450.txt | grep "^<title>" \
    | sed "s/[[:punct:]]//g" \
    | sed "s/<title>/<query>\n/g" \
    | sed "/^<query>/!s/$/\n</query>/g" \
    | sed "/^</!s/^//g" \
    | sed "s/^//g">query_title.txt ; \
sed -i "1i<parameters>" query_title.txt ; \
echo "</parameters>">>query_title.txt

```

• OKAPI TF-IDF

- OKAPI model for long query formula:

$$\sum_{t \in q} \log \frac{N}{df_t} \cdot \frac{(k_1 + 1)tf_{td}}{k_1((1 - b) + b \times (L_d/L_{ave})) + tf_{td}} \cdot \frac{(k_3 + 1)tf_{tq}}{k_3 + tf_{tq}}$$

- In order to transform OKAPI model to OKAPI TF-IDF, we have to set $k_3 = 0$
- Therefore, to use OKAPI TF-IDF on `IndexRunQuery`, we have to add
`<baseline>okapi,k1:2.0,b:0.75,k3:0.0</baseline>` in the parameter file

• Language Model with Laplace smoothing

- Laplace smoothing formula:

$$\begin{aligned} \hat{P}(w_i | c) &= \frac{\text{count}(w_i, c) + 1}{\sum_{w \in V} (\text{count}(w, c) + 1)} \\ &= \frac{\text{count}(w_i, c) + 1}{\left(\sum_{w \in V} \text{count}(w, c) \right) + |V|} \end{aligned}$$

- Dirichlet smoothing formula:

$$p(w|d) = \frac{c(w, d) + \mu p(w|C)}{|d| + \mu}$$

- Comparing the above two formulas, the Laplace smoothing can be regarded as a special case of the Dirichlet smoothing

- reference to : <http://www.cs.cmu.edu/~lemur/1.1/api.html>
(<http://www.cs.cmu.edu/~lemur/1.1/api.html>)

- Therefore, to use Language Model with Laplace smoothing in `IndexRunQuery`, we have to :

- Add `<rule>method:dirichlet,mu: unique terms</rule>` in the parameter file
 - assign unique terms to μ make it equals to number of unique terms in corpus
- Revise `indri-5.10/include/indri/DirichletTermScoreFunction.hpp`

- Before revised as shown in Figure 8 :

```
44 double scoreOccurrence( double occurrences, int contextSize ) {
45     double seen = ( double(occurrences) + _muInCollectionFrequency ) / ( double(contextSize) + _mu );
46     return log( seen );
47 }
```

Fig.8

- After revised as shown in Figure 9 :

```
44 double scoreOccurrence( double occurrences, int contextSize ) {
45     double seen = ( double(occurrences) + 1 ) / ( double(contextSize) + _mu );
46     return log( seen );
47 }
```

Fig.9

• Language Model with Jelinek-Mercer smoothing

- In order to Language Model with Jelinek-Mercer smoothing in `IndexRunQuery`, we have to add `<rule>method:jm,collectionLambda:0.8</rule>` in the parameter file

• Improve one of the above three IR models

- I choose OKAPI model as base model, and adjust the parameters k_1 , b , k_3
- In the text book-**Introduction to Information Retrieval** ch11, it mentioned:

In the absence of such optimization, experiments have shown reasonable values are to set **k_1 and k_3 to a value between 1.2 and 2** and **$b = 0.75$** .

- Therefore, I add `<baseline>okapi,k1:1.5,b:0.75,k3:1.5</baseline>` to the parameters file

Evaluation and Results

• Overview of how to evaluate

- How to evaluate retrieved results : `Trec_eval -q qrels.401-450 <retrieved result>`
`> eval_result`
- An example of retrieved results from `IndriRunQuery` as shown in Figure 10.

```
1 401 Q0 WT24-B21-59 1 15.0091 Exp
2 401 Q0 WT24-B21-104 2 15.0034 Exp
3 401 Q0 WT23-B13-52 3 13.8211 Exp
4 401 Q0 WT02-B14-10 4 13.4299 Exp
```

Fig.10

The format of retrieved results: query-

number>query-number Q0 document-id rank score Exp

- An example of evaluation of retrieved results from `Trec_eval` as shown in Figure 11.

```
1602 Queryid (Num): 50
1603 Total number of documents over all queries
1604 Retrieved: 48182
1605 Relevant: 2279
1606 Rel ret: 1471
1607 Interpolated Recall - Precision Averages:
1608 at 0.00 0.7007
1609 at 0.10 0.5412
1610 at 0.20 0.4190
1611 at 0.30 0.3176
1612 at 0.40 0.2657
1613 at 0.50 0.1921
1614 at 0.60 0.1464
1615 at 0.70 0.0965
1616 at 0.80 0.0493
1617 at 0.90 0.0268
1618 at 1.00 0.0061
1619 Average precision (non-interpolated) for all rel docs(averaged over queries)
1620 0.2269
1621 Precision:
1622 At 5 docs: 0.4640
1623 At 10 docs: 0.4260
1624 At 15 docs: 0.3787
1625 At 20 docs: 0.3400
1626 At 30 docs: 0.2953
1627 At 100 docs: 0.1578
1628 At 200 docs: 0.0996
1629 At 500 docs: 0.0512
1630 At 1000 docs: 0.0294
1631 R-Precision (precision after R (= num_rel for a query) docs retrieved):
1632 Exact: 0.2819
```

Fig.11

- Evaluation results of OKAPI TF-IDF

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48182
  Relevant:        2279
  Rel_ret:       1429
Interpolated Recall - Precision Averages:
  at 0.00        0.7021
  at 0.10        0.5299
  at 0.20        0.4089
  at 0.30        0.3026
  at 0.40        0.2483
  at 0.50        0.1802
  at 0.60        0.1262
  at 0.70        0.0832
  at 0.80        0.0339
  at 0.90        0.0169
  at 1.00        0.0038
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.2139
Precision:
  At 5 docs:     0.4560
  At 10 docs:    0.3960
  At 15 docs:    0.3733
  At 20 docs:    0.3380
  At 30 docs:    0.2867
  At 100 docs:   0.1504
  At 200 docs:   0.0957
  At 500 docs:   0.0493
  At 1000 docs:  0.0286
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:        0.2741

```

OKAPI TF-IDF with no stemming

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48885
  Relevant:        2279
  Rel_ret:       1648
Interpolated Recall - Precision Averages:
  at 0.00        0.6785
  at 0.10        0.5364
  at 0.20        0.4276
  at 0.30        0.3374
  at 0.40        0.2760
  at 0.50        0.2261
  at 0.60        0.1679
  at 0.70        0.1095
  at 0.80        0.0453
  at 0.90        0.0139
  at 1.00        0.0046
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.2320
Precision:
  At 5 docs:     0.3840
  At 10 docs:    0.4080
  At 15 docs:    0.3693
  At 20 docs:    0.3380
  At 30 docs:    0.2953
  At 100 docs:   0.1688
  At 200 docs:   0.1124
  At 500 docs:   0.0577
  At 1000 docs:  0.0330
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:        0.2876

```

OKAPI TF-IDF with porter stemming

- Evaluation results of Language Model with Laplace smoothing

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48182
  Relevant:        2279
  Rel_ret:       529
Interpolated Recall - Precision Averages:
  at 0.00        0.0577
  at 0.10        0.0427
  at 0.20        0.0302
  at 0.30        0.0198
  at 0.40        0.0130
  at 0.50        0.0114
  at 0.60        0.0101
  at 0.70        0.0095
  at 0.80        0.0062
  at 0.90        0.0053
  at 1.00        0.0007
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.0158
Precision:
  At 5 docs:     0.0120
  At 10 docs:    0.0220
  At 15 docs:    0.0267
  At 20 docs:    0.0300
  At 30 docs:    0.0300
  At 100 docs:   0.0226
  At 200 docs:   0.0200
  At 500 docs:   0.0147
  At 1000 docs:  0.0106
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:        0.0241

```

Laplace with no stemming

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48885
  Relevant:        2279
  Rel_ret:       526
Interpolated Recall - Precision Averages:
  at 0.00      0.0478
  at 0.10      0.0336
  at 0.20      0.0278
  at 0.30      0.0209
  at 0.40      0.0153
  at 0.50      0.0112
  at 0.60      0.0100
  at 0.70      0.0074
  at 0.80      0.0062
  at 0.90      0.0053
  at 1.00      0.0007
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.0140
Precision:
  At 5 docs:    0.0120
  At 10 docs:   0.0200
  At 15 docs:   0.0200
  At 20 docs:   0.0180
  At 30 docs:   0.0180
  At 100 docs:  0.0198
  At 200 docs:  0.0184
  At 500 docs:  0.0144
  At 1000 docs: 0.0105
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:       0.0195

```

Laplace with porter stemming

- Evaluation results of Language Model with Jelinek-Mercer smoothing

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48182
  Relevant:        2279
  Rel_ret:       1435
Interpolated Recall - Precision Averages:
  at 0.00      0.5811
  at 0.10      0.4148
  at 0.20      0.3228
  at 0.30      0.2705
  at 0.40      0.2146
  at 0.50      0.1653
  at 0.60      0.1274
  at 0.70      0.0969
  at 0.80      0.0652
  at 0.90      0.0508
  at 1.00      0.0301
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.1882
Precision:
  At 5 docs:    0.3120
  At 10 docs:   0.2980
  At 15 docs:   0.2893
  At 20 docs:   0.2720
  At 30 docs:   0.2427
  At 100 docs:  0.1430
  At 200 docs:  0.0936
  At 500 docs:  0.0496
  At 1000 docs: 0.0287
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:       0.2404

```

JM with no stemming

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48885
  Relevant:        2279
  Rel_ret:       1583
Interpolated Recall - Precision Averages:
  at 0.00      0.5604
  at 0.10      0.4163
  at 0.20      0.3300
  at 0.30      0.2634
  at 0.40      0.2148
  at 0.50      0.1816
  at 0.60      0.1432
  at 0.70      0.1123
  at 0.80      0.0730
  at 0.90      0.0496
  at 1.00      0.0279
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.1913
Precision:
  At 5 docs:    0.3040
  At 10 docs:   0.2880
  At 15 docs:   0.2707
  At 20 docs:   0.2550
  At 30 docs:   0.2273
  At 100 docs:  0.1462
  At 200 docs:  0.1004
  At 500 docs:  0.0544
  At 1000 docs: 0.0317
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:       0.2201

```

JM with porter stemming

- Evaluation results of Improved-OKAPI($k_1=1.5, k_3=1.5, b=0.75$)

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48182
  Relevant:        2279
  Rel_ret:       1471
Interpolated Recall - Precision Averages:
  at 0.00        0.7007
  at 0.10        0.5412
  at 0.20        0.4190
  at 0.30        0.3176
  at 0.40        0.2657
  at 0.50        0.1921
  at 0.60        0.1464
  at 0.70        0.0965
  at 0.80        0.0493
  at 0.90        0.0268
  at 1.00        0.0061
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.2269
Precision:
  At   5 docs:   0.4640
  At  10 docs:   0.4260
  At  15 docs:   0.3787
  At  20 docs:   0.3400
  At  30 docs:   0.2953
  At 100 docs:   0.1578
  At 200 docs:   0.0996
  At 500 docs:   0.0512
  At1000 docs:   0.0294
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:        0.2819

```

Improved OKAPI with no stemming

```

Queryid (Num):      50
Total number of documents over all queries
  Retrieved:      48885
  Relevant:        2279
  Rel_ret:       1706
Interpolated Recall - Precision Averages:
  at 0.00        0.6795
  at 0.10        0.5486
  at 0.20        0.4362
  at 0.30        0.3467
  at 0.40        0.2973
  at 0.50        0.2398
  at 0.60        0.1881
  at 0.70        0.1236
  at 0.80        0.0742
  at 0.90        0.0269
  at 1.00        0.0094
Average precision (non-interpolated) for all rel docs(averaged over queries)
0.2441
Precision:
  At   5 docs:   0.4040
  At  10 docs:   0.4160
  At  15 docs:   0.3827
  At  20 docs:   0.3430
  At  30 docs:   0.3013
  At 100 docs:   0.1754
  At 200 docs:   0.1168
  At 500 docs:   0.0591
  At1000 docs:   0.0341
R-Precision (precision after R (= num_rel for a query) docs retrieved):
  Exact:        0.2983

```

Improved OKAPI with porter stemming

Analysis

Preprocess raw evaluation results

- I used a script to summarize the all evaluation results
 - summarize table:

	Precision at 10	R-precision	Average Precision
model & stem			
no_stem_JM	0.298	0.2404	0.1882
no_stem_laplace	0.022	0.0239	0.0154
no_stem_improved_okapi	0.426	0.2819	0.2269
no_stem_tfidf	0.396	0.2741	0.2139
porter_JM	0.288	0.2201	0.1913
porter_laplace	0.018	0.0206	0.0136
porter_improved_okapi	0.416	0.2983	0.2441
porter_tfidf	0.408	0.2876	0.2320

Conclusion

- An overview of the performance of each **model + stemming method** is shown in Figure 12.
- Considering *Average precision*, *R-precision* and *Precision at 10*, **the improved-okapi model has the best performance.**




Fig.12

- Considering **average precision**, we group by each model to examine how stemming affects performance, as shown in Figure 13.
- As we can see in Figure 13, in most cases stemming improves performance, but sometimes stemming removes more information, leading to the worst results

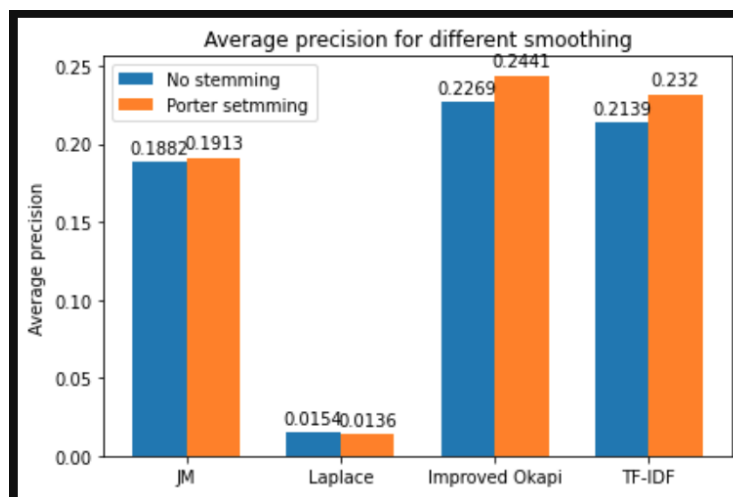


Fig.13